

Accessing Information: Recap

Registers

An x86-64 central processing unit (CPU) contains a set of 16 general-purpose registers storing 64-bit values.

Instructions can operate on data of different sizes stored in the low-order bytes of the 16 registers. Byte-level operations access the least significant byte, 16-bit operations access the least significant 2 bytes, and 32-bit operations access the least significant 4 bytes.

Different registers serve different roles in typical programs. Most unique among them is the stack pointer used to indicate the end position in the run-time stack. A set of standard programming conventions governs how the registers are to be used for managing the stack, passing function arguments, returning values from functions, and storing local (i.e., variables local to a function) and temporary data.

The following table lists these 16 registers, their names depending on the expected size of the data operands, and their roles in the standard programming convention used in this class.

Byte (8 bits)	Word (16 bits)	Long (32 bits)	Quad (64 bits)	Role
%al	%ax	%eax	%rax	Return value
%bl	%bx	%ebx	%rbx	Callee saved
%cl	%cx	%ecx	%rcx	4th argument
%dl	%dx	%edx	%rdx	3rd argument
%sil	%si	%esi	%rsi	2nd argument
%dil	%di	%edi	%rdi	1st argument
%bpl	%bp	%ebp	%rbp	Callee saved
%spl	%sp	%esp	%rsp	Stack pointer
%r8b	%r8w	%r8d	%r8	5th argument

%r9b	%r9w	%r9d	%r9	6th argument
%r10b	%r10w	%r10d	%r10	Caller saved
%r11b	%r11w	%r11d	%r11	Caller saved
%r12b	%r12w	%r12d	%r12	Callee saved
%r13b	%r13w	%r13d	%r13	Callee saved
%r14b	%r14w	%r14d	%r14	Callee saved
%r15b	%r15w	%r15d	%r15	Callee saved

Caller-saved registers are used to hold temporary values that need not be preserved across function calls. For this reason, it is the caller's responsibility to save the content of these registers (e.g., by pushing these registers onto the stack) if it needs said content after a function call.

Callee-saved registers are used to hold longer-lived values that should be preserved across function calls. When the caller makes a function call, it can expect that those registers will hold the same values after the function returns. For this reason, it is the callee's responsibility to save the content of these registers (e.g., by pushing these registers onto the stack) and to restore them before returning to the caller.

Operand Specifiers

Most instructions have one or more operands specifying the source values to use in performing an operation and the destination into which to place the result. x86-64 supports three operand types.

Type	Form	Operand value	Name
Immediate	$\$Imm$	Imm	Immediate
Register	r_a	$R[r_a]$	Register
Memory	$Imm(r_b, r_i, s)$	$M[Imm + R[r_b] + R[r_i] \cdot s]$	

With the Memory type, at least one of Imm , r_b , or r_i must be provided; if not specified the associated quantities (Imm , $R[r_b]$, $R[r_i]$ are assigned value 0). This gives rise to the following valid combinations and names.

Form	Operand value	Name
------	---------------	------

<i>Imm</i>	$M[Imm]$	Absolute
(r_b)	$M[R[r_b]]$	Indirect
$Imm(r_b)$	$M[Imm + R[r_b]]$	Base + displacement
(r_b, r_i)	$M[R[r_b] + R[r_i]]$	Indexed
$Imm(r_b, r_i)$	$M[Imm + R[r_b] + R[r_i]]$	Indexed
$(, r_i, s)$	$M[R[r_i] \cdot s]$	Scaled indexed
$Imm(, r_i, s)$	$M[Imm + R[r_i] \cdot s]$	Scaled indexed
(r_b, r_i, s)	$M[R[r_b] + R[r_i] \cdot s]$	Scaled indexed
$Imm(r_b, r_i, s)$	$M[Imm + R[r_b] + R[r_i] \cdot s]$	Scaled indexed

The scaling factor s , if specified, must be 1, 2, 4, or 8.

Data Movement Instructions

There are four simple move instructions: *movb*, *movw*, *movl*, and *movq*. They primarily differ in the amount of data that they transfer: 1, 2, 4, and 8 bytes, respectively. The source can be an Immediate, a Register, or Memory. The destination can be a Register or Memory. The sole restriction is that both cannot be Memory. When specifying a Register as a source or destination, its size designation (e.g., *%al*, *%ax*, *%eax*, or *%rax*) must match the type of the move instruction (i.e., match the suffix *b*, *w*, *l*, or *q* of the instruction). In most cases, the move instruction only updates the specific bytes of the register or memory destination. The only exception is when the destination is a 32-bit register; in that case the move instruction will also set the high-order 4 bytes of the register to 0.

The *movq* instruction can only have Immediate source operands that can be represented as 32-bit two's complement numbers. The value is then signed extended to produce the 64-bit value for the destination. The *movabsq* instruction can have an arbitrary Immediate value as its source operand and can only have a Register as a destination.

There are two groups of move instructions that deal with moving a value from a smaller size source to a larger size destination.

Instruction	Effect	Description
<i>movzbw</i>	$R_{16} \leftarrow \text{ZeroExtend}(S_8)$	Move zero-extended byte to word
<i>movzbl</i>	$R_{32} \leftarrow \text{ZeroExtend}(S_8)$	Move zero-extended byte to double word
<i>movzbq</i>	$R_{64} \leftarrow \text{ZeroExtend}(S_8)$	Move zero-extended byte to quad word
<i>movzwl</i>	$R_{32} \leftarrow \text{ZeroExtend}(S_{16})$	Move zero-extended word to double word

movzwb	$R_{64} \leftarrow \text{ZeroExtend}(S_{16})$	Move zero-extended double word to quad word
movsbw	$R_{16} \leftarrow \text{SignExtend}(S_8)$	Move sign-extended byte to word
movsbl	$R_{32} \leftarrow \text{SignExtend}(S_8)$	Move sign-extended byte to double word
movsbq	$R_{64} \leftarrow \text{SignExtend}(S_8)$	Move sign-extended byte to quad word
movswl	$R_{32} \leftarrow \text{SignExtend}(S_{16})$	Move sign-extended word to double word
movswq	$R_{64} \leftarrow \text{SignExtend}(S_{16})$	Move sign-extended word to quad word
movslq	$R_{64} \leftarrow \text{SignExtend}(S_{32})$	Move sign-extended double word to quad word
cvtq	$\%rax \leftarrow \text{SignExtend}(\%eax)$	Move sign-extended %eax to %rax

The source operand must be either a Register or Memory; destination must be a Register. When specifying a Register as a source, its size designation must be appropriate for the instruction (i.e., must match the second to last suffix letter). Likewise the size designation of the destination Register must match the last suffix letter. In most cases the instruction performs the expected operation. The exception is when the destination is a 32-bit register; in that case will also set the high-order 4 bytes of the register to 0 (even in the sign-extending cases!).

The *cvtq* instruction has no operands; it always uses *%eax* as its source and *%rax* as its destination. It performs the same operation as *movslq %eax, %rax* but it has a more compact encoding.

Pushing and Popping Stack Data

The final two data movement operations push data onto and pop data from the program stack.

Instruction	Effect	Description
pushq S	$R[\%rsp] \leftarrow R[\%rsp] - 8$ $M[R[\%rsp]] \leftarrow S$	Push quad word
popq D	$D \leftarrow M[R[\%rsp]]$ $R[\%rsp] \leftarrow R[\%rsp] + 8$	Pop quad word

The source operand for the *pushq* instruction can be an Immediate, a Register, or Memory; while the destination operand for the *popq* instruction can be a Register or Memory.

